

Scheduling Algorithms

- **FCFS (First Comes First Served)**
 - Non-Preemptive and Simplest CPU-scheduling algorithm
 - Process requesting CPU first is allocated CPU first
- **Shortest-Job-First Scheduling (Non-Preemptive)**
 - Assign CPU to the process with smallest next CPU burst
 - Next CPU bursts of 2 processes same, FCFS scheduling
- **Shortest-Remaining-Time-First (SRTF)**
 - A preemptive version of SJF algorithm
 - If the next CPU burst of the newly arrived process is shorter than what is left of the currently executing process, it will preempt the currently executing process

▶

1

1

Priority Scheduling

- A priority is associated with each process
- CPU is allocated to highest priority process
 - Equal-priority processes are scheduled in FCFS order
- Priorities are generally indicated by some fixed range of numbers, such as 0 to 7 or 0 to 1,023
- External (Static)
 - non-preemptive
- Internal (Dynamic)
 - preemptive

▶


2

2


Priority Scheduling

- Static Priority-non-preemptive
 - Higher number → Higher Priority

Process	Arrival	Burst	Priority
P1	1	6	1
P2	3	3	2
P3	6	5	3
P4	7	2	2



1 7



3

Priority Scheduling

- Static Priority-non-preemptive
 - Higher number → Higher Priority

→
→
→

Process	Arrival	Burst	Priority
P1	1	6	1
P2	3	3	2
P3	6	5	3
P4	7	2	2



1 7 12 15



4

Priority Scheduling

- Static Priority-non-preemptive
 - Higher number → Higher Priority

Process	Arrival	Burst	Priority
P1	1	6	1
P2	3	3	2
P3	6	5	3
P4	7	2	2




5

Priority Scheduling

- Dynamic Priority-preemptive
 - Higher number → Higher Priority

Process	Arrival	Burst	Priority
P1	1	6	1
P2	3	3	2
P3	6	5	3
P4	7	2	4


Ready	Running	→	Ready	Running
P1(6)	-		P2(3)	P1(4)




6

Priority Scheduling

- Dynamic Priority-preemptive
 - Higher number → Higher Priority




Process	Arrival	Burst	Priority
P1	1	6	1
P2	3	3	2
P3	6	5	3
P4	7	2	4




Priority Scheduling

- Dynamic Priority-preemptive
 - Higher number → Higher Priority




Process	Arrival	Burst	Priority
P1	1	6	1
P2	3	3	2
P3	6	5	3
P4	7	2	4




Priority Scheduling

- Dynamic Priority-preemptive
 - Higher number → Higher Priority




Process	Arrival	Burst	Priority
P1	1	6	1
P2	3	3	2
P3	6	5	3
P4	7	2	4



9

Priority Scheduling

- Dynamic Priority-preemptive
 - Higher number → Higher Priority



Process	Arrival	Burst	Priority
P1	1	6	1
P2	3	3	2
P3	6	5	3
P4	7	2	4



10

Priority Scheduling


- Dynamic Priority-preemptive
 - Higher number → Higher Priority

Ready	Running
P1(4)	P3(4)

→

Ready	Running
-	P1(4)

Process	Arrival	Burst	Priority
P1	1	6	1
P2	3	3	2
P3	6	5	3
P4	7	2	4




11

Priority Scheduling

- Dynamic Priority-preemptive
 - Higher number → Higher Priority

Ready	Running
-	-

Process	Arrival	Burst	Priority
P1	1	6	1
P2	3	3	2
P3	6	5	3
P4	7	2	4



12

Round Robin Scheduling

- Designed especially for timesharing systems
- FCFS + Preemption (to context switch processes)
- Time quantum/slice defined (10 to 100 milliseconds)
- Ready queue is a circular queue
- Time slice > time required for a typical interaction

13

Round Robin Scheduling

Ready	Running	→	Ready	Running
P1(7)	-		-	P1(4)

Process	Arrival Time	Burst Time
P1	1	7
P2	2	4
P3	3	2
P4	5	5

Time slot = 3 units

P1

1 4

14


14

Round Robin Scheduling

Ready	Running	→	Ready	Running
P2(4)	P1(4)		P3(2)	P2(4)
P3(2)			P1(4)	

Process	Arrival Time	Burst Time
P1	1	7
P2	2	4
P3	3	2
P4	5	5

Time slot = 3 units




15

Round Robin Scheduling

Ready	Running	→	Ready	Running
P3(2)	P2(1)		P1(4)	P3(2)
P1(4)			P4(5)	
P4(5)			P2(1)	



Process	Arrival Time	Burst Time
P1	1	7
P2	2	4
P3	3	2
P4	5	5

Time slot = 3 units





16

Round Robin Scheduling



17

Round Robin Scheduling




18

Round Robin Scheduling

Ready	Running	→	Ready	Running
P1(1)	P2(1)		P4(2)	P1(1)
P4(2)				

Process	Arrival Time	Burst Time
P1	1	7
P2	2	4
P3	3	2
P4	5	5

Time slot = 3 units




19

Round Robin Scheduling

Ready	Running	→	Ready	Running
P4(2)	P1(1)		-	P4(2)



Process	Arrival Time	Burst Time
P1	1	7
P2	2	4
P3	3	2
P4	5	5

Time slot = 3 units



20

Round Robin Scheduling



21

Round Robin Scheduling

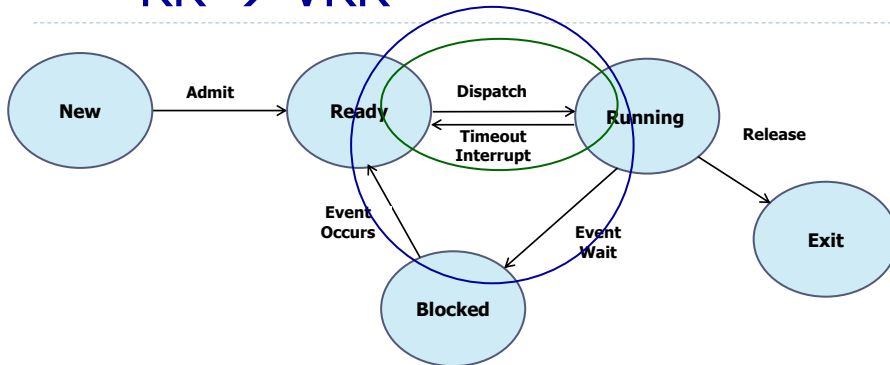
- Effective in a general-purpose time-sharing system
- **Unfair scheduling for I/O-bound processes**
- To avoid this unfairness, refinement to Round Robin is explored

Virtual Round Robin (VRR)

22

22

RR → VRR



- New → Ready → Running ↔ Ready
- Running → Wait (Blocked on IO)
- Compute Bound Vs IO Bound Processes
- VRR (RR + Auxiliary Ready Queue for IO bound processes)

23

23

Virtual Round Robin (VRR)

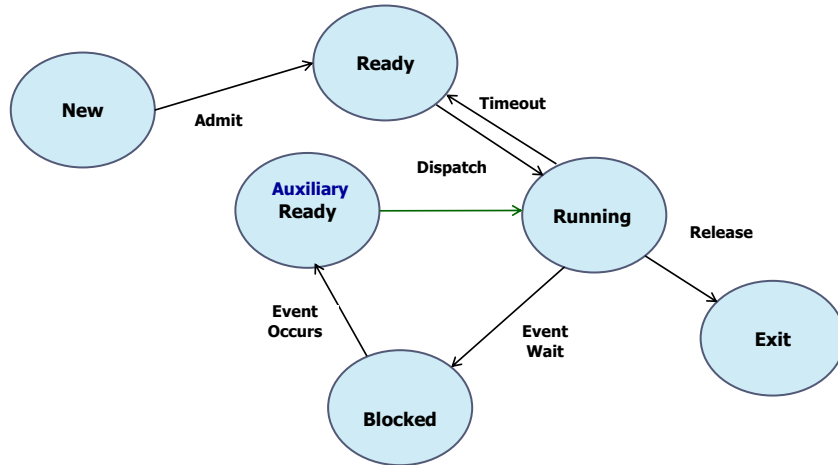
Auxiliary Queue (FCFS)

- Processes moved to Auxiliary Queue after being released from an I/O block
- When a dispatching CPU decision is to be made, processes in the Auxiliary queue get preference over main ready queue
 - A process dispatched from auxiliary queue, runs its remaining time of time slice last allocated when it initiates an I/O

24

24

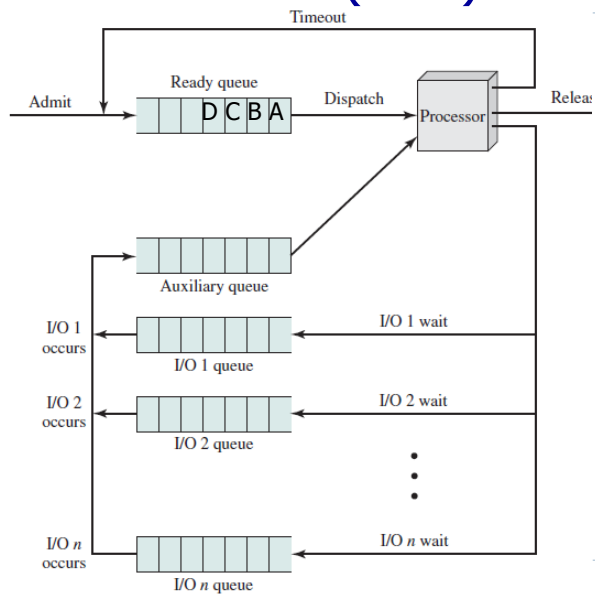
Process Model of VRR



25

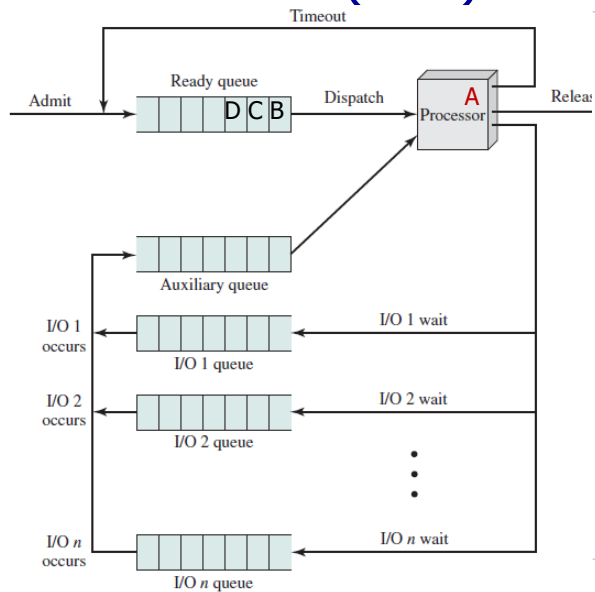
25

Virtual Round Robin (VRR)



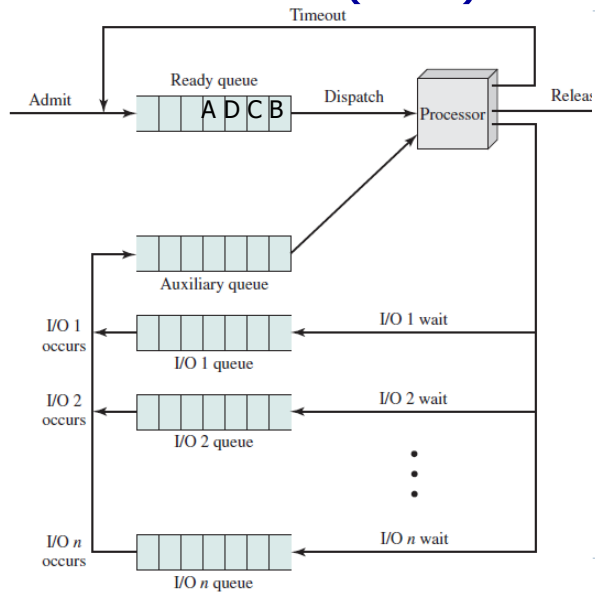
26

Virtual Round Robin (VRR)



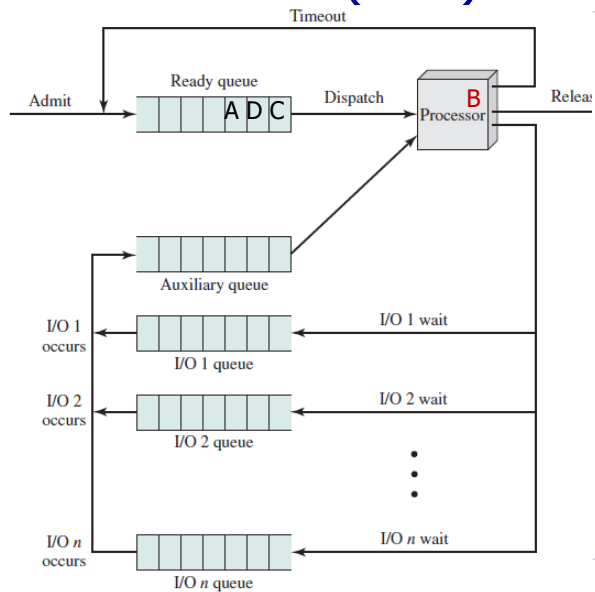
27

Virtual Round Robin (VRR)



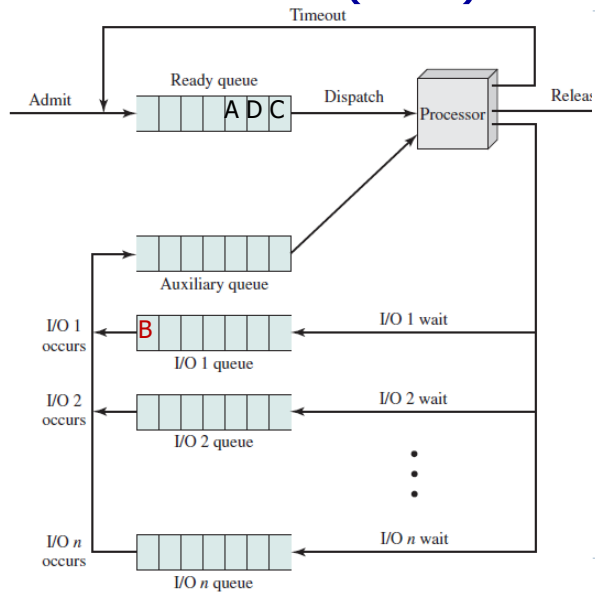
28

Virtual Round Robin (VRR)



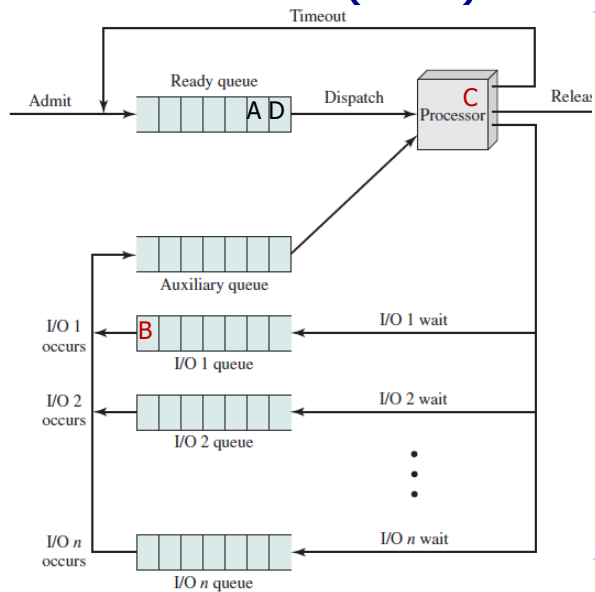
29

Virtual Round Robin (VRR)



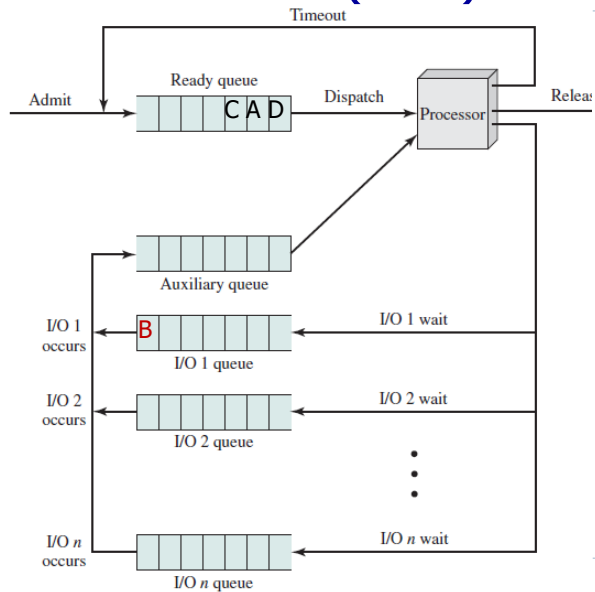
30

Virtual Round Robin (VRR)



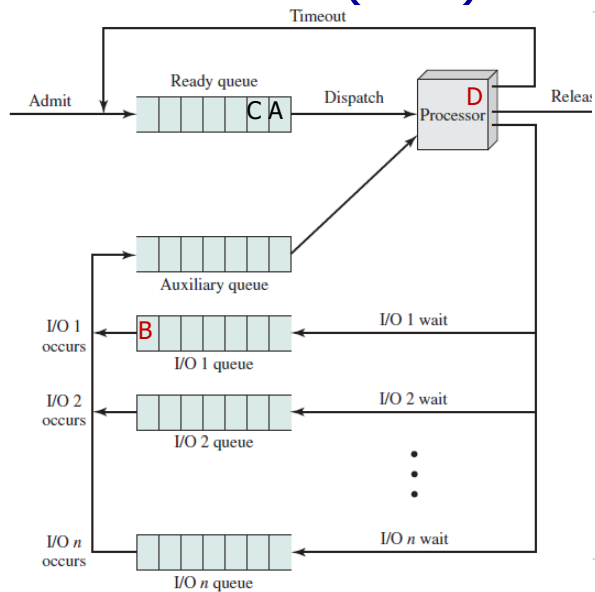
31

Virtual Round Robin (VRR)



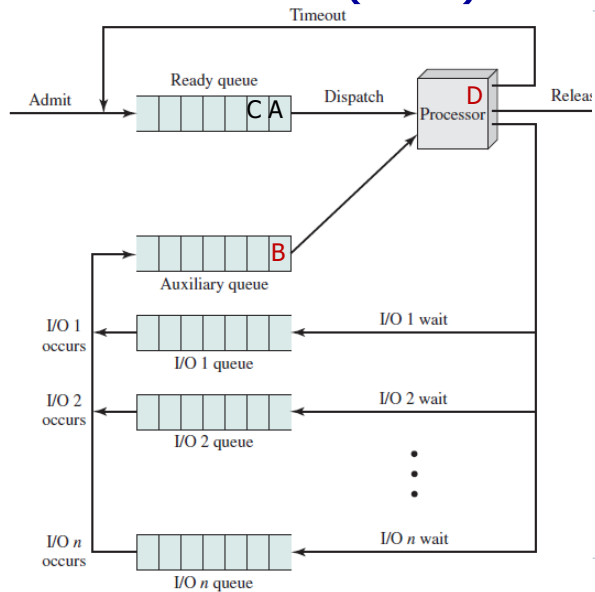
32

Virtual Round Robin (VRR)



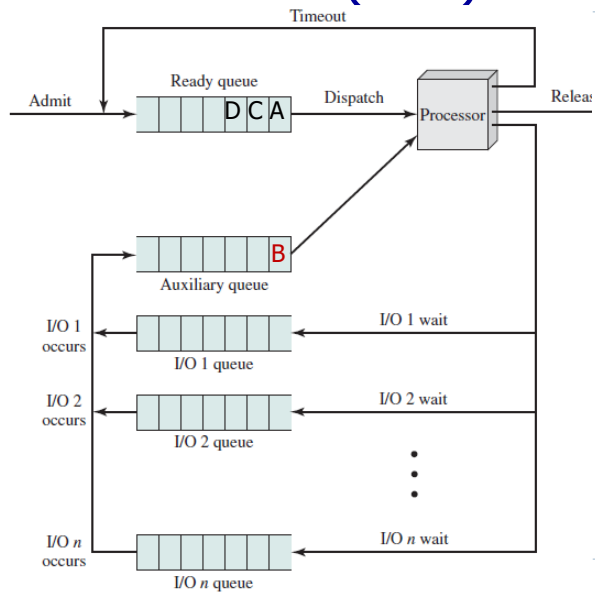
33

Virtual Round Robin (VRR)



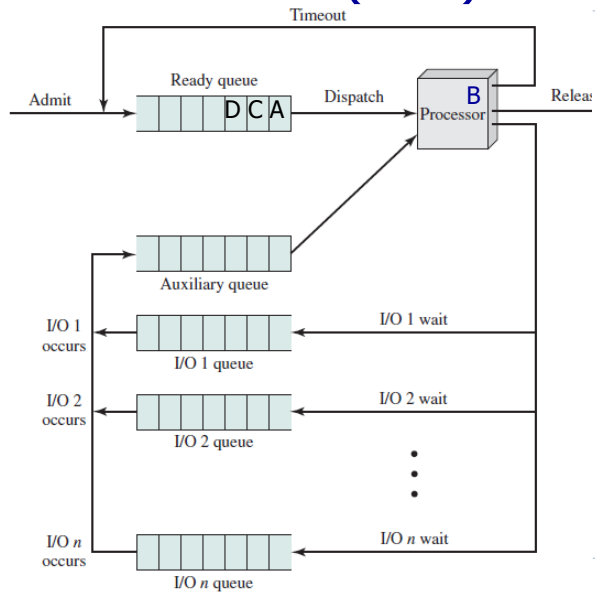
34

Virtual Round Robin (VRR)



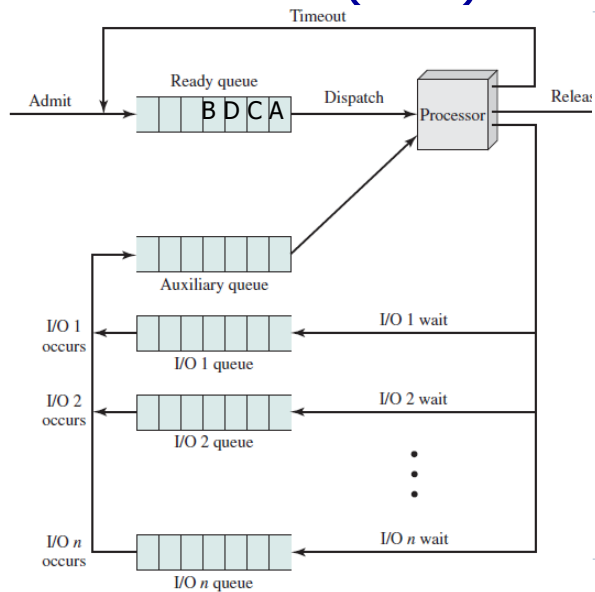
35

Virtual Round Robin (VRR)



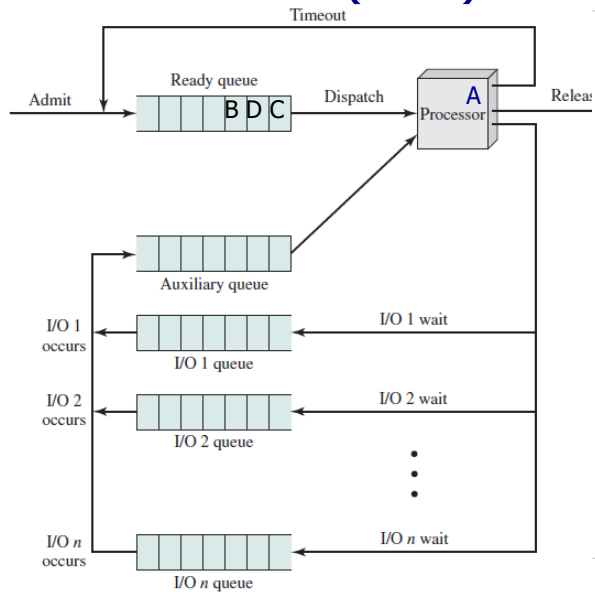
36

Virtual Round Robin (VRR)



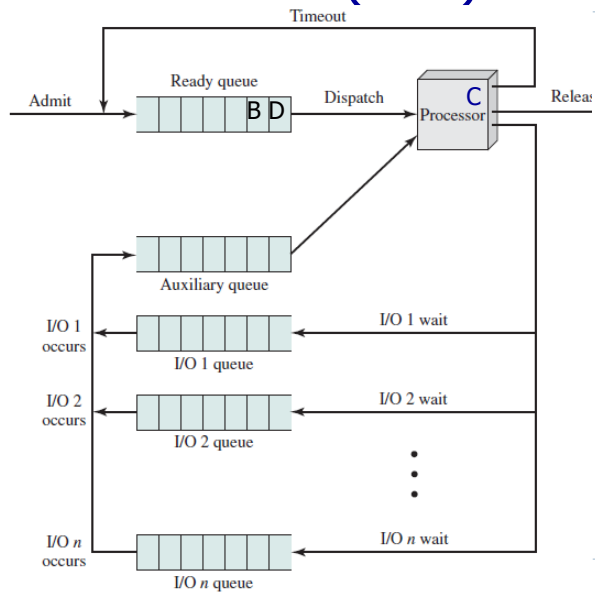
37

Virtual Round Robin (VRR)



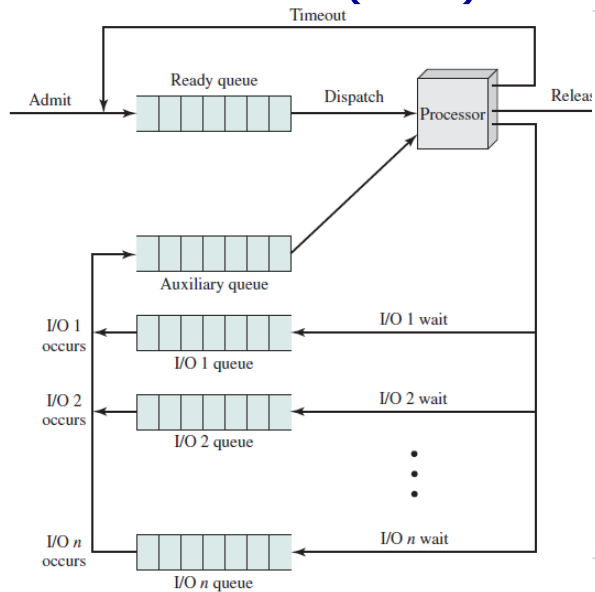
38

Virtual Round Robin (VRR)




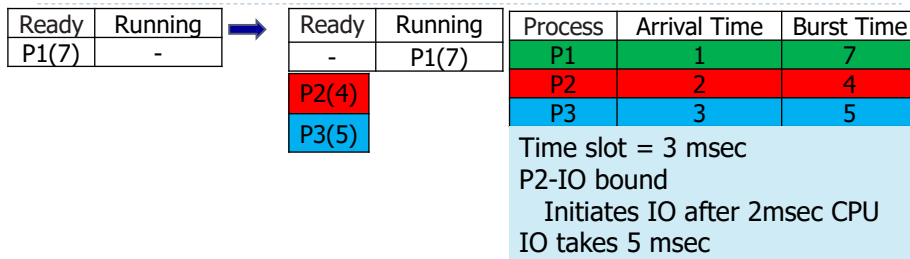
39

Virtual Round Robin (VRR)




40

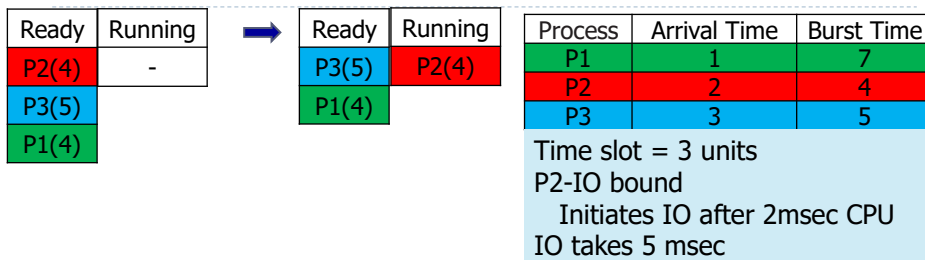
Virtual Round Robin




1 4




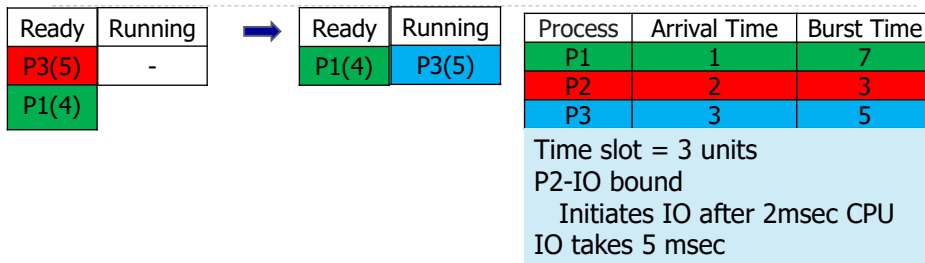
Virtual Round Robin



1 4 6



Virtual Round Robin



IO Blocked Queue


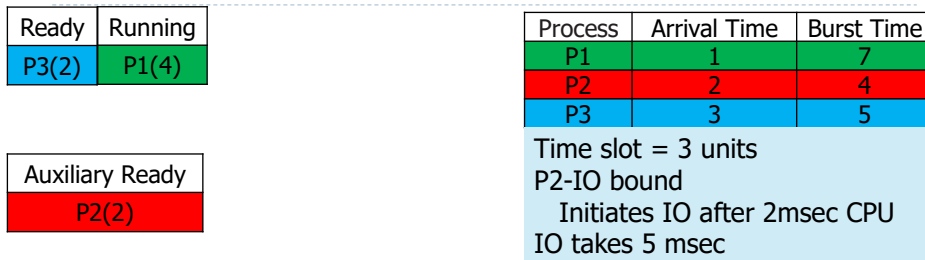


11


43

43

Virtual Round Robin



IO Blocked Queue



11

44

44


Virtual Round Robin

Ready	Running
P3(2)	P1(4)

Auxiliary Ready
P2(2)

Process	Arrival Time	Burst Time
P1	1	7
P2	2	4
P3	3	5

Time slot = 3 units
P2-I/O bound
Initiates IO after 2msec CPU
IO takes 5 msec



IO Blocked Queue



45


Virtual Round Robin

Ready	Running
P3(2)	-
P1(1)	

Auxiliary Ready
P2(2)

Process	Arrival Time	Burst Time
P1	1	7
P2	2	3
P3	3	5

Time slot = 3 units
P2-I/O bound
Initiates IO after 2msec CPU
IO takes 5 msec

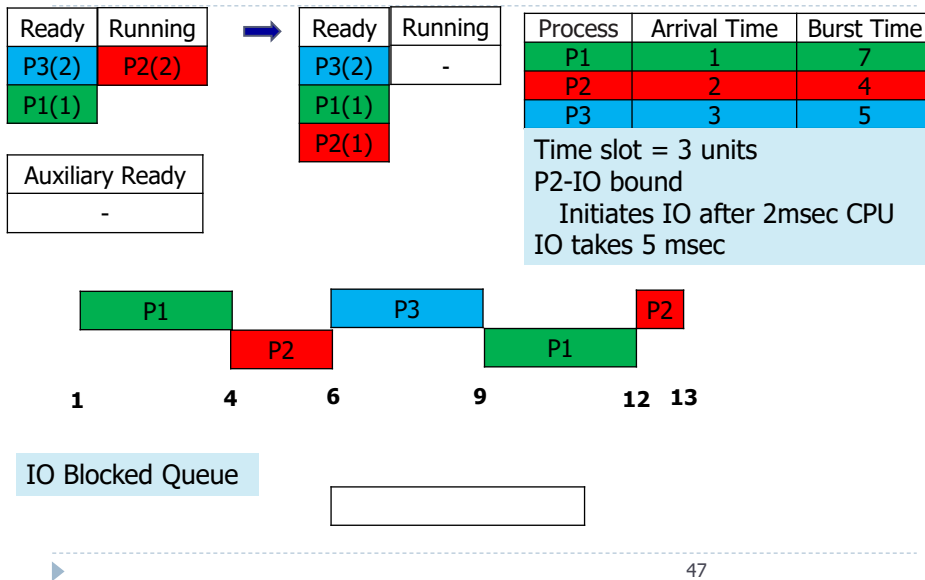


IO Blocked Queue



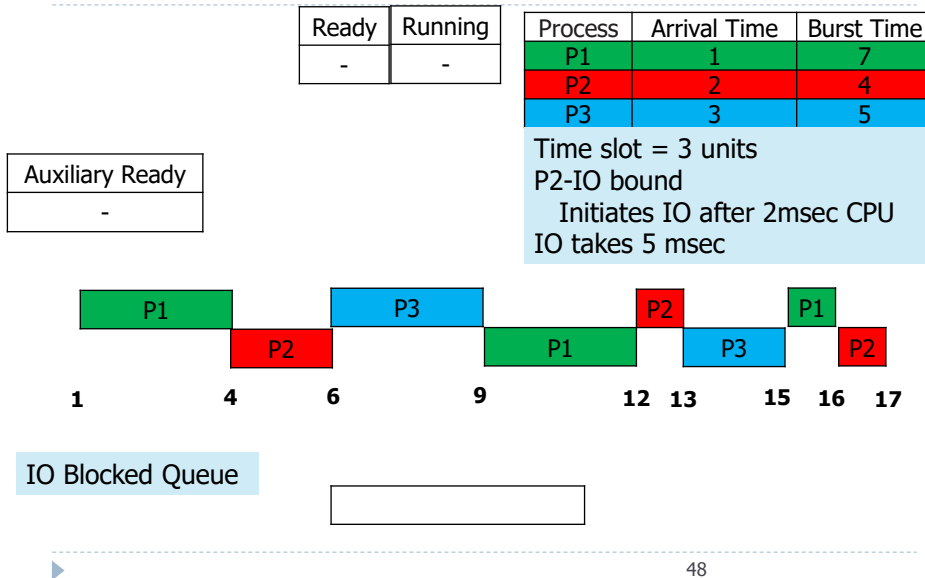
46

Virtual Round Robin



47

Virtual Round Robin



48

Multilevel Queue Scheduling

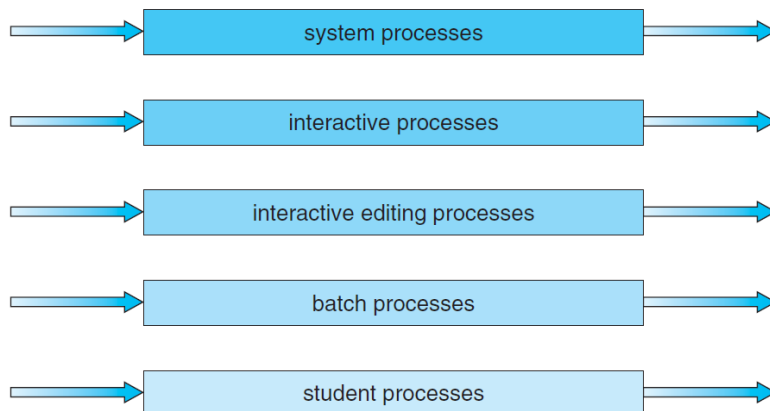
- Partitions ready queue into several separate queues
- Processes are permanently assigned to one queue
 - memory size, process priority, or process type
- Each queue has its own scheduling algorithm
 - foreground and background processes
 - foreground queue might be scheduled by RR
 - background queue is scheduled by FCFS
- Scheduling among the queues (fixed-priority preemptive scheduling)

49

49

Multilevel Queue Scheduling

highest priority



lowest priority

50

50

Multilevel Queue Scheduling

- Time-slicing among the queues
 - Each queue gets a certain portion of the CPU time
 - schedule among its various processes
- foreground–background queue example
 - foreground queue with 80 percent of the CPU time for RR scheduling among its processes
 - background queue with 20 percent of the CPU to give to its processes on an FCFS basis

▶

51

51

Multilevel Feedback Queue Scheduling

- Exploiting past behavior
- Multiple queues, each with different priority
- Allows a process to move between queues
- Idea is to separate processes according to the characteristics of their CPU bursts
- Each queue has its own scheduling algorithm
 - e.g. foreground – RR, background – FCFS
- Sometimes multiple RR priorities with quantum increasing exponentially (highest:4ms, next:8ms, next:16ms, etc.)
 - A process uses too much CPU time → move to a lower-priority queue


▶

52

52

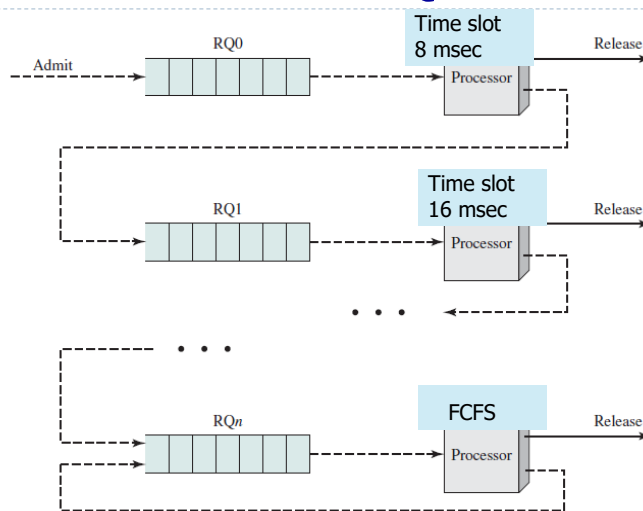
Multilevel Feedback Queue Scheduling

- Job starts in highest priority queue
- If timeslot expires, drop one level
- If timeslot doesn't expire, push up one level (or to top)



53

Multilevel Feedback Queue Scheduling



54